

Kinshuk Jagdev

Logistic Regression

Task

Build a predictive model that answers the question: “what sorts of people were more likely to survive?” using passenger data (ie name, age, gender, socio-economic class, etc).

In [1]:

```
import warnings
warnings.filterwarnings('ignore')
```

1: Data Loading

First and Foremost Step is loading the data. We will use **Pandas** library to load the data into pandas **Dataframe**.

Main Library Used: Pandas

In [2]:

```
#importing pandas library
import pandas as pd

#importing numpy library
import numpy as np
```

In [3]:

```
#Loading training dataset into dataframe

df_train = pd.read_csv('E:\\ml data\\train.csv')
```


In [4]:

```
#previewing the training data
```

```
df_train.head()
```

Out[4]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare (
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250
3	4	1	1	Futelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500



In [5]:

```
#obtaining data information
```

```
df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 891 entries, 0 to 890  
Data columns (total 12 columns):  
PassengerId    891 non-null int64  
Survived       891 non-null int64  
Pclass         891 non-null int64  
Name           891 non-null object  
Sex            891 non-null object  
Age           714 non-null float64  
SibSp          891 non-null int64  
Parch          891 non-null int64  
Ticket         891 non-null object  
Fare           891 non-null float64  
Cabin          204 non-null object  
Embarked       889 non-null object  
dtypes: float64(2), int64(5), object(5)  
memory usage: 83.7+ KB
```

Analysis: Training dataset has 891 Entries

In [6]:

```
#Loading the test data into dataframe
```

```
df_test = pd.read_csv('E:\\ml data\\test.csv')
```

In [7]:

#previewing the test data

df_test.head()

Out[7]:

	PassengerId	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Emba
0	892	3	Kelly, Mr. James	male	34.5	0	0	330911	7.8292	NaN	
1	893	3	Wilkes, Mrs. James (Ellen Needs)	female	47.0	1	0	363272	7.0000	NaN	
2	894	2	Myles, Mr. Thomas Francis	male	62.0	0	0	240276	9.6875	NaN	
3	895	3	Wirz, Mr. Albert	male	27.0	0	0	315154	8.6625	NaN	
4	896	3	Hirvonen, Mrs. Alexander (Helga E Lindqvist)	female	22.0	1	1	3101298	12.2875	NaN	

In [8]:

#obtaining test data info

df_test.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 418 entries, 0 to 417
Data columns (total 11 columns):
PassengerId    418 non-null int64
Pclass         418 non-null int64
Name           418 non-null object
Sex            418 non-null object
Age            332 non-null float64
SibSp          418 non-null int64
Parch          418 non-null int64
Ticket         418 non-null object
Fare           417 non-null float64
Cabin          91 non-null object
Embarked       418 non-null object
dtypes: float64(2), int64(4), object(5)
memory usage: 36.0+ KB
```

Analysis: Testing dataset has 418 Entries

2: Data Cleaning and Filling Null Values

We will be dropping **Name**, **Ticket** and **PassengerId** variables, as they are mostly unique and have no

significance in analysis

In [9]:

```
df_train = df_train.drop(['Name', 'Ticket', 'PassengerId'], 1)
df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 9 columns):
Survived      891 non-null int64
Pclass        891 non-null int64
Sex           891 non-null object
Age           714 non-null float64
SibSp         891 non-null int64
Parch         891 non-null int64
Fare          891 non-null float64
Cabin         204 non-null object
Embarked      889 non-null object
dtypes: float64(2), int64(4), object(3)
memory usage: 62.8+ KB
```

In [10]:

```
passenger_data = df_test['PassengerId']
df_test = df_test.drop(['Name', 'Ticket', 'PassengerId'], 1)
df_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 418 entries, 0 to 417
Data columns (total 8 columns):
Pclass        418 non-null int64
Sex           418 non-null object
Age           332 non-null float64
SibSp         418 non-null int64
Parch         418 non-null int64
Fare          417 non-null float64
Cabin         91 non-null object
Embarked      418 non-null object
dtypes: float64(2), int64(3), object(3)
memory usage: 26.2+ KB
```

Getting the count of null values

In [11]:

```
df_train.isnull().sum()
```

Out[11]:

```
Survived      0
Pclass        0
Sex            0
Age          177
SibSp         0
Parch         0
Fare          0
Cabin        687
Embarked      2
dtype: int64
```

Analysis: Age and Cabin have 177 and 687 null values, while Embarked has only 2 null values

Since **Cabin** has roughly **77%** missing data we will **drop** this column

Age column has roughly **20%** missing data, we will simply fill null values with **mean** age values

In [12]:

```
df_train = df_train.drop(['Cabin'], 1)
```

```
#viewing new data
df_train.head()
```

Out[12]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	0	3	male	22.0	1	0	7.2500	S
1	1	1	female	38.0	1	0	71.2833	C
2	1	3	female	26.0	0	0	7.9250	S
3	1	1	female	35.0	1	0	53.1000	S
4	0	3	male	35.0	0	0	8.0500	S

Filling null values in Embarked Column

This column has three stations from which people have boarded, so we will simply fill the null values with most common station, i.e. S

In [13]:

```
df_train["Embarked"].fillna(df_train["Embarked"].value_counts().idxmax(), inplace=True)
```

Filling null values in Age Column

Filling mean values

In [14]:

```
df_train['Age'].fillna(df_train['Age'].mean(skipna=True), inplace = True)
```

Analysing new data

In [15]:

```
df_train.isnull().sum()
```

Out[15]:

```
Survived    0
Pclass      0
Sex          0
Age         0
SibSp       0
Parch       0
Fare        0
Embarked    0
dtype: int64
```

Conclusion: : Data now has no null values, so we will repeat same process for testing data

In [16]:

```
df_test.isnull().sum()
```

Out[16]:

```
Pclass      0
Sex          0
Age         86
SibSp       0
Parch       0
Fare        1
Cabin      327
Embarked    0
dtype: int64
```

In [17]:

```
df_test = df_test.drop(['Cabin'], 1)
```

```
#viewing new data
df_test.head()
```

Out[17]:

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	3	male	34.5	0	0	7.8292	Q
1	3	female	47.0	1	0	7.0000	S
2	2	male	62.0	0	0	9.6875	Q
3	3	male	27.0	0	0	8.6625	S
4	3	female	22.0	1	1	12.2875	S

In [18]:

```
df_test["Embarked"].fillna(df_test['Embarked'].value_counts().idxmax(), inplace=True)
```

In [19]:

```
df_test['Age'].fillna(df_test['Age'].mean(skipna=True), inplace = True)
```

Fare Column

Fare Column has a missing value in test data, but not in train data, so we will simply replace it with mean value.

In [20]:

```
df_test['Fare'].fillna(df_test['Fare'].mean(skipna=True), inplace = True)
```

Checking new data

In [21]:

```
df_test.isnull().sum()
```

Out[21]:

```
Pclass      0
Sex          0
Age          0
SibSp        0
Parch        0
Fare         0
Embarked     0
dtype: int64
```

3: Dealing with Categorical Variables

We have SibSp and Parch variables, which tells us if an individual is travelling with family. So, we will simply create a variable, **travelling_alone** to optimize our data.

In [22]:

```
#creating travel_alone variable

df_train['travel_alone']=np.where((df_train["SibSp"]+df_train["Parch"])>0, 0, 1)
df_train.drop('SibSp', axis=1, inplace=True)
df_train.drop('Parch', axis=1, inplace=True)
```

Getting Dummy Variables

The Dummy Variables enable us to use a single regression equation to represent multiple groups

In [23]:

#creating dummy variables bor "Sex" column

```
dum1 = pd.get_dummies(df_train, columns=['Sex'], prefix = 'Sex')
df_train = pd.concat([dum1], 1)
df_train.head()
```

Out[23]:

	Survived	Pclass	Age	Fare	Embarked	travel_alone	Sex_female	Sex_male
0	0	3	22.0	7.2500	S	0	0	1
1	1	1	38.0	71.2833	C	0	1	0
2	1	3	26.0	7.9250	S	1	1	0
3	1	1	35.0	53.1000	S	0	1	0
4	0	3	35.0	8.0500	S	1	0	1

In [24]:

#dropping extra variable

```
df_train = df_train.drop(['Sex_female'], 1)
```

In [25]:

#checking the data

```
df_train.head()
```

Out[25]:

	Survived	Pclass	Age	Fare	Embarked	travel_alone	Sex_male
0	0	3	22.0	7.2500	S	0	1
1	1	1	38.0	71.2833	C	0	0
2	1	3	26.0	7.9250	S	1	0
3	1	1	35.0	53.1000	S	0	0
4	0	3	35.0	8.0500	S	1	1

In [26]:

```
#dummy variables for class and embarked variables
```

```
dm2=pd.get_dummies(df_train, columns=["Pclass","Embarked"])
dm2.head()
```

Out[26]:

	Survived	Age	Fare	travel_alone	Sex_male	Pclass_1	Pclass_2	Pclass_3	Embarked_C
0	0	22.0	7.2500	0	1	0	0	1	0
1	1	38.0	71.2833	0	0	1	0	0	1
2	1	26.0	7.9250	1	0	0	0	1	0
3	1	35.0	53.1000	0	0	1	0	0	0
4	0	35.0	8.0500	1	1	0	0	1	0

In [27]:

```
df_train = pd.concat([dm2], 1)
df_train.head()
```

Out[27]:

	Survived	Age	Fare	travel_alone	Sex_male	Pclass_1	Pclass_2	Pclass_3	Embarked_C
0	0	22.0	7.2500	0	1	0	0	1	0
1	1	38.0	71.2833	0	0	1	0	0	1
2	1	26.0	7.9250	1	0	0	0	1	0
3	1	35.0	53.1000	0	0	1	0	0	0
4	0	35.0	8.0500	1	1	0	0	1	0

Repeating same steps for test data

In [28]:

```
df_test['travel_alone']=np.where((df_test["SibSp"]+df_test["Parch"])>0, 0, 1)
df_test.drop('SibSp', axis=1, inplace=True)
df_test.drop('Parch', axis=1, inplace=True)
```

Dummy Variables for Test data

In [29]:

```
dum1 = pd.get_dummies(df_test, columns=['Sex'], prefix = 'Sex')
df_test = pd.concat([dum1], 1)
df_test.head()
```

Out[29]:

	Pclass	Age	Fare	Embarked	travel_alone	Sex_female	Sex_male
0	3	34.5	7.8292	Q	1	0	1
1	3	47.0	7.0000	S	0	1	0
2	2	62.0	9.6875	Q	1	0	1
3	3	27.0	8.6625	S	1	0	1
4	3	22.0	12.2875	S	0	1	0

In [30]:

```
df_test = df_test.drop(['Sex_female'], 1)
```

In [31]:

```
df_test.head()
```

Out[31]:

	Pclass	Age	Fare	Embarked	travel_alone	Sex_male
0	3	34.5	7.8292	Q	1	1
1	3	47.0	7.0000	S	0	0
2	2	62.0	9.6875	Q	1	1
3	3	27.0	8.6625	S	1	1
4	3	22.0	12.2875	S	0	0

In [32]:

```
dm2=pd.get_dummies(df_test, columns=["Pclass","Embarked"])
dm2.head()
```

Out[32]:

	Age	Fare	travel_alone	Sex_male	Pclass_1	Pclass_2	Pclass_3	Embarked_C	Embarke
0	34.5	7.8292	1	1	0	0	1	0	
1	47.0	7.0000	0	0	0	0	1	0	
2	62.0	9.6875	1	1	0	1	0	0	
3	27.0	8.6625	1	1	0	0	1	0	
4	22.0	12.2875	0	0	0	0	1	0	

In [33]:

```
df_test = pd.concat([dm2], 1)
df_test.head()
```

Out[33]:

	Age	Fare	travel_alone	Sex_male	Pclass_1	Pclass_2	Pclass_3	Embarked_C	Embarke
0	34.5	7.8292	1	1	0	0	1	0	
1	47.0	7.0000	0	0	0	0	1	0	
2	62.0	9.6875	1	1	0	1	0	0	
3	27.0	8.6625	1	1	0	0	1	0	
4	22.0	12.2875	0	0	0	0	1	0	

Both test and training data now are in desired shape, now it is time to start with outlier treatment

4: Univariate Analysis

Since our data is in good condition now, it is time to start understanding data with Univariate Analysis. Univariate Analysis is the simplest analysis method.

Main Library Used: Seaborn

In [34]:

```
#importing necessary libraries for visualuzation of data

import matplotlib.pyplot as plt
import seaborn as sns
```


Survival Based on Passenger Class

In [35]:

```
sns.countplot("Pclass_1", hue="Survived", data = df_train)
```

Out[35]:

<matplotlib.axes._subplots.AxesSubplot at 0x24021b93e48>




In [36]:

```
sns.countplot("Pclass_2", hue="Survived", data = df_train)
```

Out[36]:

<matplotlib.axes._subplots.AxesSubplot at 0x24021eb0248>

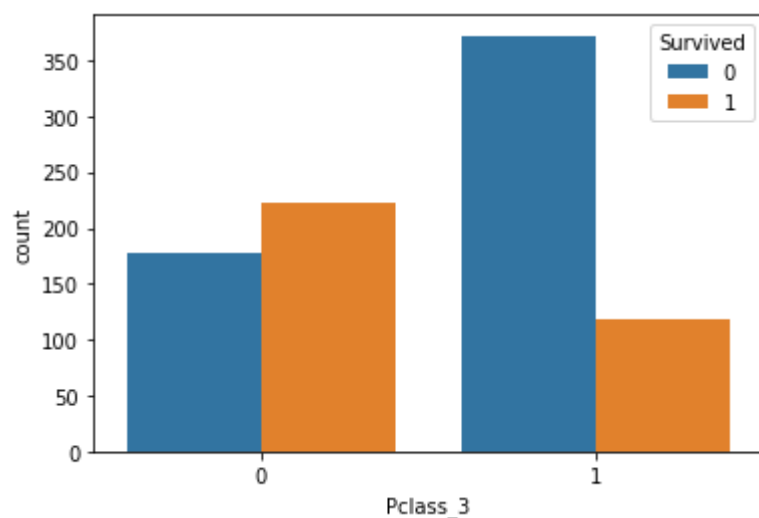


In [37]:

```
sns.countplot("Pclass_3", hue="Survived", data = df_train)
```

Out[37]:

<matplotlib.axes._subplots.AxesSubplot at 0x24021f32208>



Survival Percentages in Passenger Class

In [38]:

```
df_train[["Pclass_1", "Survived"]].groupby(['Pclass_1'], as_index=False).mean().sort_values
```

Out[38]:

	Pclass_1	Survived
1	1	0.629630
0	0	0.305185

In [39]:

```
df_train[["Pclass_2", "Survived"]].groupby(['Pclass_2'], as_index=False).mean().sort_values
```

Out[39]:

	Pclass_2	Survived
1	1	0.472826
0	0	0.360679

In [40]:

```
df_train[["Pclass_3", "Survived"]].groupby(['Pclass_3'], as_index=False).mean().sort_values
```

Out[40]:

	Pclass_3	Survived
0	0	0.557500
1	1	0.242363

In [41]:

```
print('*'*50)
print("Class-1 Survival Rate: 62%")
print("Class-2 Survival Rate: 47%")
print("Class-3 Survival Rate: 24%")
print('*'*50)
```

```
*****
Class-1 Survival Rate: 62%
Class-2 Survival Rate: 47%
Class-3 Survival Rate: 24%
*****
```

Analysis: Class-1 Passengers have very high survival percentage while Class-3 Passengers have least survival rate

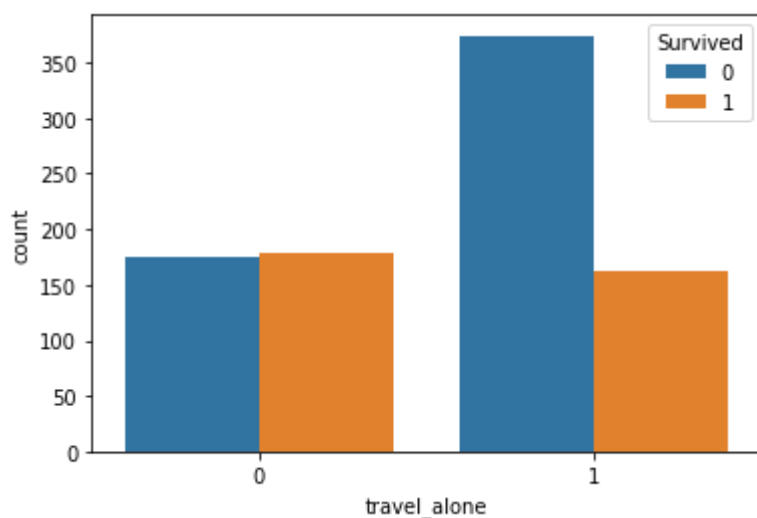
Survival Based on If one is travelling alone

In [42]:

```
sns.countplot('travel_alone', hue='Survived', data=df_train)
```

Out[42]:

<matplotlib.axes._subplots.AxesSubplot at 0x24021fbf348>



Survival Percentages in travel_alone

In [43]:

```
df_train[["travel_alone", "Survived"]].groupby(['travel_alone'], as_index=False).mean().sort
```

Out[43]:

	travel_alone	Survived
0	0	0.505650
1	1	0.303538

Analysis: Those travelling with families have had higher survival rate


Survival Based on Gender

In [44]:

```
sns.countplot('Sex_male', hue='Survived', data=df_train)
```

Out[44]:

<matplotlib.axes._subplots.AxesSubplot at 0x2402202d388>



Survival Percentages in Gender

In [45]:

```
df_train[["Sex_male", "Survived"]].groupby(['Sex_male'], as_index=False).mean().sort_values
```

Out[45]:

Sex_male	Survived
0	0.742038
1	0.188908

Analysis: Females have higher survival rate than Males


Survival Based on Age

In [46]:

```
sns.countplot('Age', hue='Survived', data=df_train)
```

Out[46]:

<matplotlib.axes._subplots.AxesSubplot at 0x24022095888>




Survival Based on Fare

In [47]:

```
sns.countplot('Fare', hue='Survived', data=df_train)
```

Out[47]:

<matplotlib.axes._subplots.AxesSubplot at 0x24022399288>




In [49]:

```
sns.countplot("Embarked_C", hue="Survived", data = df_train)
```

Out[49]:

<matplotlib.axes._subplots.AxesSubplot at 0x24022ba7d08>




In [50]:

```
sns.countplot("Embarked_Q", hue="Survived", data = df_train)
```

Out[50]:

<matplotlib.axes._subplots.AxesSubplot at 0x24022d1df88>



Survival Percentages in Different Boarding Locations

In [51]:

```
df_train[["Embarked_C", "Survived"]].groupby(['Embarked_C'], as_index=False).mean().sort_va
```

Out[51]:

	Embarked_C	Survived
1	1	0.553571
0	0	0.344398

In [52]:

```
df_train[["Embarked_S", "Survived"]].groupby(['Embarked_S'], as_index=False).mean().sort_va
```

Out[52]:

	Embarked_S	Survived
0	0	0.502041
1	1	0.339009

In [53]:

```
df_train[["Embarked_Q", "Survived"]].groupby(['Embarked_Q'], as_index=False).mean().sort_va
```

Out[53]:

	Embarked_Q	Survived
1	1	0.389610
0	0	0.383292

In [54]:

```
print('*'*50)
print("Port-C Survival Rate: 55%")
print("Port-S Survival Rate: 34%")
print("Port-Q Survival Rate: 39%")
print('*'*50)
```

```
*****
Port-C Survival Rate: 55%
Port-S Survival Rate: 34%
Port-Q Survival Rate: 39%
*****
```

Analysis: Port-C has highest survival rate while Port-S has least survival rate

Adding final categorical variable, for age, as it is important column for analysis

In [55]:

```
df_train['Minor']=np.where(df_train['Age']<=16, 1, 0)
df_test['Minor']=np.where(df_test['Age']<=16, 1, 0)
```


5: Logistic Regression Model Building

Since we are now familiar with our data, now we can finally perform Logistic Regression **Main Library Used:** Sklearn

Initially Checking The Correlation:

In [56]:

```
plt.figure(figsize = (8,6))
sns.heatmap(df_train.corr(), annot = True)
plt.show()
```



Building the model

Splitting X and Y training data

In [57]:

```
#using sklearn for Logistic Regression
from sklearn.linear_model import LogisticRegression

#getting RFE for feature selection
from sklearn.feature_selection import RFE
```

In [58]:

```
#getting initial variables
vars = ["Age", "Fare", "travel_alone", "Pclass_1", "Pclass_2", "Embarked_C", "Embarked_S", "Sex_ma

#splitting into X and Y training datasets
X_train = df_train[vars]
y_train = df_train['Survived']
```

Feature Selection Using RFE

In [59]:

```
#creating logistic regresssion model
md = LogisticRegression()

#using rfe for feature selection
rfe = RFE(md, 8)
rfe = rfe.fit(X_train, y_train)
list(X_train.columns[rfe.support_])
```

Out[59]:

```
['Age',
 'travel_alone',
 'Pclass_1',
 'Pclass_2',
 'Embarked_C',
 'Embarked_S',
 'Sex_male',
 'Minor']
```

Finally, we have our variables ready, so we can move forward with model building

In [60]:

```
#updating training and testing datasets

vars = ['Age', 'travel_alone', 'Pclass_1', 'Pclass_2', 'Embarked_C', 'Embarked_S', 'Sex_male',
X_train = df_train[vars]
X_test = df_test[vars]
X_train.columns
```

Out[60]:

```
Index(['Age', 'travel_alone', 'Pclass_1', 'Pclass_2', 'Embarked_C',
      'Embarked_S', 'Sex_male', 'Minor'],
      dtype='object')
```

In [61]:

```
#fitting final model
logreg = LogisticRegression()
logreg.fit(X_train, y_train)
```

Out[61]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='warn', n_jobs=None, penalty='l2',
                    random_state=None, solver='warn', tol=0.0001, verbose=0,
                    warm_start=False)
```

6: Testing the Model

In [62]:

```
#Testing the data and saving predictions

y_pred = logreg.predict(X_test)
```

Getting VIF values

In [63]:

```
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

In [64]:

```
vif = pd.DataFrame()
vif['vars'] = X_train.columns
vif['VIF'] = [variance_inflation_factor(X_train.values, i) for i in range(X_train.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[64]:

	vars	VIF
0	Age	6.84
5	Embarked_S	6.04
6	Sex_male	3.21
1	travel_alone	2.90
4	Embarked_C	2.38
2	Pclass_1	1.88
3	Pclass_2	1.48
7	Minor	1.43

Analysis : VIF values are normal, so we can continue with this model

Creating Dataframe for Submission

In [65]:

```
#passenger ids from test data
l1 = passenger_data.tolist()

#creating dataframe
#logreg_submission = pd.DataFrame({"PassengerId": l1, "Survived": y_pred})

#dataframe to csv conversion
#logreg_submission.to_csv('2ndlogreg_submission.csv', index=False)
```

7: Post Submission Analysis:

Model Score: 0.7727

The Model Score recieved is actually good, and the best the we could achieve at the moment. Although, we acknowledge that there is scope for improvement.

Thank you